MATRIX MULTIPLICATION
STEP BY STEP SPEED UP

NISANTH M P, 2023 JAN 27

Agenda

- To talk about agenda
- Motivation
- Definition of matrix multiplication
- Step by step optimisations demos
- Reasons for the gains at each step
- More things to try out

Not part of agenda: MMA

Motivation

- Performance optimisation is a vast field spanning numerous hardware architectures and workloads from various domains
- Workloads from many domains have matrix algebra at its core - so fast BLAS is a vast and highly active research field
- Most BLAS operations could be formulated such that they boil down to GEMM - GEneral Matrix Multiplication
- Modern BLAS libraries are structured such that majority of the code is generic for all hardware architectures, with only tuneable parameters and GEMM kernels in hand assembly (or intrinsics) specific for a given HW architecture
- We will see how highly optimised generic code for GEMM could be intuitively developed - leaving the hand assembly and specific micro architecture performance analysis alone for the moment

The Definition

If $\mathbf{A}$ is an $m \times n$ matrix and $\mathbf{B}$ is an $n \times p$ matrix,

$$\mathbf{A} = \begin{pmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{pmatrix}, \quad \mathbf{B} = \begin{pmatrix} b_{11} & b_{12} & \cdots & b_{1p} \\ b_{21} & b_{22} & \cdots & b_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ b_{n1} & b_{n2} & \cdots & b_{np} \end{pmatrix}$$

the *matrix product* $\mathbf{C} = \mathbf{AB}$ (denoted without multiplication signs or dots) is defined to be the $m \times p$ matrix

$$\mathbf{C} = \begin{pmatrix} c_{11} & c_{12} & \cdots & c_{1p} \\ c_{21} & c_{22} & \cdots & c_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ c_{m1} & c_{m2} & \cdots & c_{mp} \end{pmatrix}$$

such that

$$c_{ij} = a_{i1}b_{1j} + a_{i2}b_{2j} + \cdots + a_{in}b_{nj} = \sum_{k=1}^{n} a_{ik}b_{kj},$$

for $i = 1, ..., m$ and $j = 1, ..., p$.

Reference

Step 01: Definition in Code

```c
for(int i = 0; i < n; ++i)
    for(int j = 0; j < n; ++j)
        for(int k = 0; k < n; ++k)
        {
            C[i][j] += A[i][k] * B[k][j];
        }
```

n = 2048, GCC 13.0, No Optimisation: ~50 seconds

n = 2048, GCC 13.0, O3 Optimisation: ~17 seconds
(ie. ~3x improvement with compiler optimisations - no change in code)

System Details

Step 02: Swap Indices

```
for(int i = 0; i < n; ++i)
    for(int k = 0; k < n; ++k)
        for(int j = 0; j < n; ++j)
        {
            C[i][j] += A[i][k] * B[k][j];
        }
```

Step 01: n = 2048, GCC 13.0, O3 Optimisation: ~17 seconds

Step 02: n = 2048, GCC 13.0, O3: ~2.32 seconds
(ie. ~7x improvement with single line of change in code)

Reason?   System Details

Step 03: Loop Unroll k

```
for(int i = 0; i < n; ++i)
    for(int k = 0; k < n; ++k)
        for(int j = 0; j < n; j += 16)
        {
            C(i,j) += A(i,k) * B(k,j);
            C(i,j+1) += A(i,k) * B(k,j+1);
            C(i,j+2) += A(i,k) * B(k,j+2);
            C(i,j+3) += A(i,k) * B(k,j+3);
            C(i,j+4) += A(i,k) * B(k,j+4);
            C(i,j+5) += A(i,k) * B(k,j+5);
            C(i,j+6) += A(i,k) * B(k,j+6);
            C(i,j+7) += A(i,k) * B(k,j+7);
            C(i,j+8) += A(i,k) * B(k,j+8);
            C(i,j+9) += A(i,k) * B(k,j+9);
            C(i,j+10) += A(i,k) * B(k,j+10);
            C(i,j+11) += A(i,k) * B(k,j+11);
            C(i,j+12) += A(i,k) * B(k,j+12);
            C(i,j+13) += A(i,k) * B(k,j+13);
            C(i,j+14) += A(i,k) * B(k,j+14);
            C(i,j+15) += A(i,k) * B(k,j+15);
        }
```

Step 02: n = 2048, GCC 13.0, O3: ~2.32 seconds

Step 03: n = 2048, GCC 13.0, O3: ~XXXX seconds

System Details

Step 03: Loop Unroll j

```
for(int i = 0; i < n; ++i)
    for(int k = 0; k < n; ++k)
        for(int j = 0; j < n; j += 16)
        {
            C(i,j) += A(i,k) * B(k,j);
            C(i,j+1) += A(i,k) * B(k,j+1);
            C(i,j+2) += A(i,k) * B(k,j+2);
            C(i,j+3) += A(i,k) * B(k,j+3);
            C(i,j+4) += A(i,k) * B(k,j+4);
            C(i,j+5) += A(i,k) * B(k,j+5);
            C(i,j+6) += A(i,k) * B(k,j+6);
            C(i,j+7) += A(i,k) * B(k,j+7);
            C(i,j+8) += A(i,k) * B(k,j+8);
            C(i,j+9) += A(i,k) * B(k,j+9);
            C(i,j+10) += A(i,k) * B(k,j+10);
            C(i,j+11) += A(i,k) * B(k,j+11);
            C(i,j+12) += A(i,k) * B(k,j+12);
            C(i,j+13) += A(i,k) * B(k,j+13);
            C(i,j+14) += A(i,k) * B(k,j+14);
            C(i,j+15) += A(i,k) * B(k,j+15);
        }
```

Step 02: n = 2048, GCC 13.0, O3: ~2.32 seconds

Step 03: n = 2048, GCC 13.0, O3: ~2.30 seconds

NO CHANGE!!

Reason?

System Details

Step 04: Loop Unroll i

```
for(int i = 0; i < n; i += 4)
    for(int k = 0; k < n; ++k)
        for(int j = 0; j < n; j += 16)
        {
            C(i,j) += A(i,k) * B(k,j);
            C(i,j+1) += A(i,k) * B(k,j+1);
            ;
            ;
            C(i,j+15) += A(i,k) * B(k,j+15);

            C(i+1,j) += A(i+1,k) * B(k,j);
            C(i+1,j+1) += A(i+1,k) * B(k,j+1);
            ;
            ;
            C(i+1,j+15) += A(i+1,k) * B(k,j+15);

            C(i+2,j) += A(i+2,k) * B(k,j);
            C(i+2,j+1) += A(i+2,k) * B(k,j+1);
            ;
            ;
            C(i+2,j+15) += A(i+2,k) * B(k,j+15);

            C(i+3,j) += A(i+3,k) * B(k,j);
            C(i+3,j+1) += A(i+3,k) * B(k,j+1);
            ;
            ;
            C(i+3,j+15) += A(i+3,k) * B(k,j+15);
        }
```

Step 03: n = 2048, GCC 13.0, O3: ~2.30 seconds

Step 04: n = 2048, GCC 13.0, O3: ~1.20 seconds

~2x improvement

Reason?

System Details

Step 05: Blocking

```c
int mb = 128, nb = 128;
for (int nk = 0; nk < n; nk += nb)
    for (int mi = 0; mi < n; mi += mb) {
        double *a = &AS[mi][nk];
        double *b = &BS[nk][0];
        double *c = &CS[mi][0];

        // Matrix multiplication
        for(int i = 0; i < mb; i += 4)
            for(int k = 0; k < nb; ++k)
                for(int j = 0; j < n; j += 16)
                {
                    C(i,j) += A(i,k) * B(k,j);
                    ;
                    C(i,j+15) += A(i,k) * B(k,j+15);

                    C(i+1,j) += A(i+1,k) * B(k,j);
                    ;
                    C(i+1,j+15) += A(i+1,k) * B(k,j+15);

                    C(i+2,j) += A(i+2,k) * B(k,j);
                    ;
                    C(i+2,j+15) += A(i+2,k) * B(k,j+15);

                    C(i+3,j) += A(i+3,k) * B(k,j);
                    ;
                    C(i+3,j+15) += A(i+3,k) * B(k,j+15);
                }
    }
```

Step 04: n = 2048, GCC 13.0, O3:
~1.20 seconds

Step 05: n = 2048, GCC 13.0, O3:
~0.70 seconds

~2x improvement

Reason?

System Details

Step 06: Multithreading

```c
int mb = 128, nb = 128;
#pragma omp parallel for
for (int nk = 0; nk < n; nk += nb)
    for (int mi = 0; mi < n; mi += mb) {
        double *a = &AS[mi][nk];
        double *b = &BS[nk][0];
        double *c = &CS[mi][0];

        // Matrix multiplication
        for(int i = 0; i < mb; i += 4)
            for(int k = 0; k < nb; ++k)
                for(int j = 0; j < n; j += 16)
                {
                    C(i,j) += A(i,k) * B(k,j);
                    ;
                    C(i,j+15) += A(i,k) * B(k,j+15);

                    C(i+1,j) += A(i+1,k) * B(k,j);
                    ;
                    C(i+1,j+15) += A(i+1,k) * B(k,j+15);

                    C(i+2,j) += A(i+2,k) * B(k,j);
                    ;
                    C(i+2,j+15) += A(i+2,k) * B(k,j+15);

                    C(i+3,j) += A(i+3,k) * B(k,j);
                    ;
                    C(i+3,j+15) += A(i+3,k) * B(k,j+15);
                }
    }
```

Step 05: n = 2048, GCC 13.0, O3:
~0.7 seconds

Step 06: n = 2048, GCC 13.0, O3:
~0.18 seconds

~3.5x improvement

System Details

System Details

system: IBM,9080-HEX

processor: PowerPC,POWER10

CPU clock: 3450MHz

CPU(s): 384

L1d cache: 32KB

L1i cache: 48KB

L2 cache: 1024KB

L3 cache: 4096KB

memory: 477GiB System memory

OS: Linux

Reasons

Step 01: Increased cache/register hits and xxxx

Reasons

Step 02: Increased cache/register hits and compiler optimised SIMD

Step 03: Compiler already did this ^ loop unrolling and SIMD optimisation

Back to step: <u>03</u>

Reasons

Step 02: Increased cache/register hits and compiler optimised SIMD
Step 03: Compiler already did this ^ loop unrolling and SIMD optimisation
Step 04: Increased cache/register hits

Back to step: 04

Reasons

Step 02: Increased cache/register hits and compiler optimised SIMD
Step 03: Compiler already did this ^ loop unrolling and SIMD optimisation
Step 04: Increased cache/register hits
Step 05: Increased cache hits due to smaller blocks fitting in L2 cache

More Things to Try

07: Hand assembly optimisations to include special vector instructions etc
    - specific for ISAs

08: Experimenting with different loop unrolling factors - depends on register sizes
    and register file sizes

09: Experimenting with different block sizes - depends on L2 and L3 cache sizes

10: Packing blocks of matrices A & B - to increase cache hits

11: Experimenting with different thread nesting strategies to maximise parallelism -
    core scaling

12: Experimenting with different thread pinning strategies to optimise for NUMA
    scaling

References

Have borrowed some ideas from BLIS' GEMM speed up wiki - https://github.com/flame/how-to-optimize-gemm

Code and other documents available at: https://github.com/nisanthmp/matmul-speedup

Q&A

Thank You